



Bachelor's project

Parallel Bit-Counting

for Approximate Similarity Searching

Rasmus Hag Løvstad, pgq596@alumni.ku.dk

Supervisor: Mikkel Thorup

Handed in: January 16, 2022

Abstract

This project regards an analysis of recent advances in solving the *Approximate Similarity Search Problem* with respect to the *Jaccard Similarity* $J(A, B) = \frac{A \cap B}{A \cup B}$, for some subsets A, B of some universe U . Given a corpus \mathcal{F} consisting of n subsets of U , two constants $0 < j_2 < j_1 < 1$ as well as a set $Q \subseteq U$ find a set $A \in \mathcal{F}$ such that $J(A, Q) \geq j_2$ if there exists a set $B \in \mathcal{F}$ where $J(B, Q) \geq j_1$.

First, existing solutions to this problem will be described, such as the *Locality Sensitive Hashing Framework* (LSH) by Gionis, Indyk, and Motwani [3], the seminal *MinHash* algorithm by Broder [1] and the *Fast Similarity Sketching* algorithm by Dahlgaard, Knudsen, and Thorup [7]. Then the focus will be shifted to the most recent advances by Knudsen [11], who presents a parallel bit-counting algorithm that supposedly can make the query time sub-linear to the amount of sets in \mathcal{F} .

The efficiency of this bit-counting algorithm will then be investigated by evaluating the algorithm in terms of time complexity, proof of correctness and an empirical comparison to simpler methods. The findings include a theoretical time advantage to using parallel bit counting compared to a simple, linear time algorithm, but empirical results show that this advantage is hard to achieve in practice. Furthermore, reflections are made on how these results might scale on more specialized hardware, wherein the Approximate Similarity Search Problem may find the most use.

Contents

1	Introduction	4
2	Theory	6
2.1	Problem Definition	6
2.2	Trivial Solution	6
2.3	Locality Sensitive Hashing Framework	7
2.4	MinHash	7
2.4.1	Implementation with the LSH framework	8
2.5	Later Improvements	10
3	Parallel Bit Counting	12
3.1	Divide-and-Conquer	12
3.1.1	Example	14
3.2	Parallelism	14
3.2.1	Example	16
3.3	Discrepancies	17
3.4	Time Complexity Analysis	18
4	Empirical Evaluation	19
4.1	Methods	19
4.1.1	Benchmarking	19
4.2	Implementation	20
4.2.1	Technology	20
4.2.2	Design	20
4.2.3	Correctness	21

4.2.4	Benchmarking	21
4.3	Results	21
4.4	Discussion	23
5	Conclusion	25
	References	26
	Mathematical Notation	27
	Appendices	28
A	Machine Specifications	28
B	Code	28

1 Introduction

The *Approximate Similarity Search Problem* regards efficiently finding a set $A \in \mathcal{F}$ from some corpus \mathcal{F} consisting of n subsets of some universe U that is approximately similar to a query set Q in regards to the *Jaccard Similarity* metric $J(A, Q) = \frac{|A \cap Q|}{|A \cup Q|}$ [7], [11]. Approximate means that given two variables $0 < j_2 < j_1 < 1$, the algorithm should efficiently find some set $A \in \mathcal{F}$ such that $J(A, Q) \geq j_2$ if there exists some set $B \in \mathcal{F}$ where $J(B, Q) \geq j_1$. Practical applications includes searching through large corpora of high-dimensional text documents like plagiarism-detection or website duplication checking among others [9]. The main bottleneck in this problem is the *curse of dimensionality*. A trivial algorithm can solve this problem in $O(n|Q| \max_{A \in \mathcal{F}} |A|)$ time (where n is the amount of elements in the corpus \mathcal{F}), but algorithms with a query-time proportional to the dimensionality of the corpus $\max_{A \in \mathcal{F}} |A|$ scale poorly when working with high-dimensional data sets. Text documents are especially bad in this regard since they often are encoded using *w-shingles* (w contiguous words) which Li, Shrivastava, Moore, *et al.* [6] shows easily can reach a dimensionality upwards of $|Q| = 2^{83}$ using just 5-shingles.

The classic solution to this problem is the *MinHash* algorithm presented by Broder [1] to perform website duplication checking for the AltaVista search engine. It pre-processes the data once using hashing to perform effective querying in $O(n^\rho \log(n)|Q|)$ time (where $\rho = \frac{\log_2(1/j_1)}{\log_2(1/j_2)}$), a significant improvement independent of the dimensionality of the corpus. Many improvements have since been presented to both improve pre-processing time, query time and space efficiency. Notable mentions includes (but are not limited to) the use of *tensoring* [4], *b-bit minwise hashing* [5], *fast similarity sketching* [7]. Applications of these techniques lead to efficient querying in $O(n^\rho + |Q|)$ time, with a constant error probability. If one wishes to achieve an even better error probability such as $\varepsilon = o(1)$, it is standard practice within the field to use $O(\log_2(1/\varepsilon))$ independent data structures and return the best result, resulting in a query time of $O(\log_2(1/\varepsilon)(n^\rho + |Q|))$.

Recent advances by Knudsen [11] show that it is possible to achieve an even better query time by sampling these data structures from one large sketch, leading to a query time of $O(n^\rho \log_2(1/\varepsilon) + |Q|)$.

Each of the data structures will return a list of candidate sets, of which false positives needs to be eliminated from. To solve this, Knudsen [11] presents an efficient filtering scheme can choose exactly 1 set from each data structure as a candidate. Then, the final result can be picked by eliminating data structures that are highly probable to return a bad candidate. This final step requires computing the cardinality of a list of bit-strings (i.e. amount of bits set per bit-string) very efficiently, for which an algorithm is presented by Knudsen [11] that brings the query time down to $O((\frac{n \log_2 w}{w})^\rho \log(1/\varepsilon) + |Q|)$ where w is the word size by working with multiple words packed with bit-strings in parallel.

The main focus of this project is to analyze, prove, implement and evaluate this parallel bit-counting technique. The analysis will be based on the original paper [11], but with some modifications to resolve some of the issues with the original algorithm. This will also include a pseudo-code implementation of the algorithm since the original paper only describes it through recurrences. This leads to a proof of correctness that slightly differs from the one presented in the paper as well as a time complexity analysis.

The theoretical analysis will be backed up by a real-life implementation that will be

benchmarked to investigate the benefits of parallelism in practice. At last, reflections on the results and methods will be made to back up eventual conclusions.

2 Theory

The *Similarity Search Problem* has many variations and uses. The general idea is to find some data point from a set of data points that is the most similar to some query point according to some similarity metric. The LSH framework, one of the first big breakthroughs in this field, regards points in a d -dimensional space and uses the Manhattan distance between them as a similarity metric - two points that are close together are regarded as more similar [3]. In this project, we will regard the similarity between sets and use the Jaccard similarity $J(A, B) = \frac{|A \cap B|}{|A \cup B|}$ as our similarity metric. The Jaccard similarity produces a number between 0 and 1, where a larger number indicates more similarity between the two sets. It is furthermore only defined when both A and B are non-empty, as to avoid dividing by zero.

This part will first define the problem to be considered in a concise way. This will then be used to describe existing techniques and solutions to the problem, which will set the context for the next section that will specifically focus on one of the most recent advances within the field: A parallel bit-counting algorithm that can filter out potential false matches effectively.

2.1 Problem Definition

The formal definition of the *Jaccard Similarity Search Problem* is as follows: Given a family \mathcal{F} of n sets from some universe U and a query set $Q \subset U$, pre-process \mathcal{F} to efficiently find the set $A \in \mathcal{F}$ such that $A = \arg \max_{A \in \mathcal{F}} J(A, Q)$.

In many practical applications, it may be enough to only consider the *Approximate Jaccard Similarity Search Problem*: Given the setting above and two constants $0 < j_2 < j_1 < 1$, pre-process \mathcal{F} such that one can efficiently find some set $A \in \mathcal{F}$ such that $J(A, Q) \geq j_2$ if there exists some set $B \in \mathcal{F}$ where $J(B, Q) \geq j_1$. Note that it is possible for $A = B$. The intuition to this is that we can dial in how precise we expect our algorithm to be able to solve this problem by fine-tuning j_2 and j_1 . We no longer expect the algorithm to return the *best* solution, just one that is "good enough" (if a good solution exists).

Most existing methods to solve this problem depend on the relationship between j_1 and j_2 . This relationship is often described as $\rho = \frac{\log_2(1/j_1)}{\log_2(1/j_2)}$. Note that ρ shrinks as j_1 approaches 1. The query time of most solutions use ρ as a exponent of n , which means that the smaller the value of ρ is, the less the size of the corpus influences the running time.

2.2 Trivial Solution

The most obvious solution is to compare each element in each set $A \in \mathcal{F}$ with each element in Q to calculate the size of the intersection $|A \cap Q|$ and union $|A \cup Q|$. This solution will always give the correct answer, and query in $O(n|Q| \max_{A \in \mathcal{F}} |A|)$ time trivially, but suffers from the *curse of dimensionality*: The running time is proportional to the dimensionality of the corpus, which can be a big bottleneck in terms of performance. When working with high-dimensional data sets like often seen in text processing, this can have huge consequences. Li, Shrivastava, Moore, *et al.* [6] has shown that it is easy to reach a dimensionality upwards of $|A| = 2^{83}$ when using *5-shingles* (5 contiguous words) of the 10,000 most common English words. This

algorithm is guaranteed to return the set $S = \arg \max_{A \in \mathcal{F}} J(Q, A)$ on every query. The next algorithm can query much faster than this by relaxing this condition: If we allow pre-processing the data and relax the requirement of finding an exact match by considering the approximate similarity search problem instead, we can create so-called sketches of the data set before querying. This sketching strategy can drastically reduce the query time.

2.3 Locality Sensitive Hashing Framework

One of the first big contributions to this problem is the Locality Sensitive Hashing Framework by Indyk and Motwani [2]. This framework aims to solve the Approximate Similarity Search problem for any kind of similarity space, including Jaccard similarity, using some appropriate (s_1, s_2, r_1, r_2) -sensitive family of hash functions.

Definition 1. A family of hash functions $\mathcal{H} = \{h: X \rightarrow U\}$ is (s_1, s_2, r_1, r_2) -sensitive for some set X and some universe U with a similarity function $S: X \times X \rightarrow [0; 1]$ if these conditions apply for any $x, y \in X$ and any $h \in \mathcal{H}$:

- If $S(x, y) \geq s_1$, then $\Pr[h(x) = h(y)] \geq r_1$
- If $S(x, y) \leq s_2$, then $\Pr[h(x) = h(y)] \leq r_2$

First, we define a family of functions $\mathcal{G} = \{g: X \rightarrow U^k\}$ such that for each $g \in \mathcal{G}$, $g(p) = \langle h_0(p), \dots, h_{k-1}(p) \rangle$ for $h_i \in \mathcal{H}$, where $k \in \mathbb{Z}^+$ is to be determined later. We can now regard $g(p)$ as the index of a “bucket”, where we can store the value p . We do this for each of the functions $g_0, \dots, g_{l-1} \in \mathcal{G}$ for each data point $p \in X$ and some value $l \in \mathbb{Z}^+$ to be determined later. If a collision happens, we store only one of the colliding points, chosen at random.

To query, look up the corresponding point for each of the buckets $g_0(q), \dots, g_{l-1}(q)$. This yields the points $P = \{p_0, \dots, p_{t-1}\}$ where $t \leq l$ is the amount of points found. For each $p_i \in P$, if $S(p_i, q) \geq j_2$, then return it. If no points fulfill this constraint, return false.

Part of implementing this framework for a specific similarity metric (e.g. Jaccard Similarity) requires choosing a suitable value for k and l . Finding the right parameters can in this case result in getting a constant error probability, as we will show later.

2.4 MinHash

The MinHash algorithm is one of the classic examples of a locality sensitive hash family in regards to the Jaccard similarity.

It is based on the following theorem:

Theorem 1. Let h be a universal hash function from $V \cup W \rightarrow [M]$ for some $M \in \mathbb{Z}^+$ where $W \cup V \subseteq U$ for some universe U . Let $H(W) = \min_{w \in W} h(w)$ for some h picked uniformly at random. Then

$$\Pr[H(W) = H(V)] = \frac{|V \cap W|}{|V \cup W|} = J(V, W)$$

Proof. If you sample a random element $x \in V \cup W$, then you can find the probability that $x \in V \cap W$ like so:

$$Pr[x \in V \cap W] = \frac{|V \cap W|}{|V \cup W|} = J(V, W)$$

When we calculate $H(W)$ and $H(V)$, the smallest value of the two will be $H(W \cup V)$. Since we use hash functions, this is essentially equivalent of sampling a random value from $W \cup V$. Assuming no collisions, this value will be in the intersection $W \cap V$ if and only if $H(W)$ and $H(V)$ have picked the same element. This means that the probability of $H(W) = H(V)$ is the same as the probability of picking an element of the intersection from the union, which is $\frac{|V \cap W|}{|V \cup W|} = J(V, W)$. \square

The MinHash algorithm works like so: First, pre-process each set $A \in \mathcal{F}$ by hashing each element in A using k universal hash functions picked uniformly at random, saving the smallest value for each hash function as a sketch $S(A) = \langle H_0(A), \dots, H_{k-1}(A) \rangle$. To query a set Q , calculate its sketch $S(Q)$ and then calculate

$$d(S(A), S(Q)) = \frac{1}{k} \sum_{i \in [k]} [H_i(A) = H_i(Q)]$$

We can then derive the expected value of this to be the Jaccard Similarity:

$$\mathbb{E}[d(S(A), S(Q))] = \mathbb{E}\left[\frac{1}{k} \sum_{i \in [k]} [H_i(A) = H_i(Q)]\right]$$

Using linearity of expectation:

$$= \frac{1}{k} \sum_{i \in [k]} \mathbb{E}[[H_i(A) = H_i(Q)]]$$

Using theorem 1:

$$\begin{aligned} &= \frac{1}{k} \sum_{i \in [k]} J(A, Q) \\ &= J(A, Q) \end{aligned}$$

2.4.1 Implementation with the LSH framework

For a function $H \in \mathcal{H}$ from some family \mathcal{H} defined as $H(W) = \min_{w \in W} h(w)$ as in theorem 1, we know that $Pr[H(A) = H(B)] = J(A, B)$ as per theorem 1. This must mean that \mathcal{H} is (j_1, j_2, j_1, j_2) -sensitive for some values $0 < j_2 < j_1 < 1$ as per definition 1 and theorem 1:

- If $J(W, V) \geq j_1$, then $Pr[H(W) = H(V)] \geq j_1$
- If $J(W, V) \leq j_2$, then $Pr[H(W) = H(V)] \leq j_2$.

We can then consider a function $g \in \mathcal{G}$ from some family \mathcal{G} defined as $g(W) = \langle H_0(W), \dots, H_{k-1}(W) \rangle$ where functions $H_0 \dots H_{k-1}$ are picked uniformly at random

from \mathcal{H} . An evaluation of any $g \in \mathcal{G}$ will be called a “sketch” from now on. Since the functions $H_0 \dots H_{k-1}$ are independent, we can use Bayes’ theorem to derive that:

$$\begin{aligned} \Pr[g(W) = g(V)] &= \Pr\left[\bigcap_{H \in \mathcal{G}} H(W) = H(V)\right] \\ &= \prod_{H \in \mathcal{G}} \Pr[H(W) = H(V)] \\ &= \prod_{H \in \mathcal{G}} J(W, V) = J(W, V)^k \end{aligned}$$

This means that \mathcal{G} is (j_1, j_2, j_1^k, j_2^k) -sensitive for some $k \in \mathbb{Z}^+$ since $J(W, V)$ and j_1, j_2 are non-negative:

- If $J(W, V) \geq j_1$, then $J(W, V)^k \geq j_1^k$, which means that $\Pr[g(W) = g(V)] \geq j_1^k$
- If $J(W, V) \leq j_2$, then $J(W, V)^k \leq j_2^k$, which means that $\Pr[g(W) = g(V)] \leq j_2^k$

By choosing $k = \lceil \frac{\log_2(n)}{\log_2(1/j_2)} \rceil$, we can derive that \mathcal{G} now is $(j_1, j_2, 1/n^\rho, 1/n)$ -sensitive where $\rho = \frac{\log_2 1/j_1}{\log_2 1/j_2}$.

$$\begin{aligned} j_1^k &= j_1^{\lceil \frac{\log_2(n)}{\log_2(1/j_2)} \rceil} \geq j_1^{\frac{\log_2(n)}{\log_2(1/j_2)}} \\ &= j_1^{\frac{\log_2(n)}{\log_2(1/j_2)} \cdot \frac{\log_2(1/j_1)}{\log_2(1/j_1)}} \\ &= j_1^{\frac{\log_2(1/j_1)}{\log_2(1/j_2)} \cdot \frac{\log_2(n)}{\log_2(j_1)}} \\ &= j_1^{-\rho \frac{\log_2(n)}{\log_2(j_1)}} \\ &= j_1^{-\rho \log_{j_1}(n)} \\ &= \frac{1}{n^\rho} \\ j_2^k &\leq j_2 \cdot j_2^{\frac{\log_2(n)}{\log_2(1/j_2)}} = j_2 \cdot j_2^{-\frac{\log_2(n)}{\log_2(j_2)}} = \frac{j_2}{n} \leq 1/n \end{aligned}$$

Now, imagine $g_0, \dots, g_{l-1} \in \mathcal{G}$ picked uniformly at random where $l = \lceil n^\rho \rceil$. If we calculate the sketch $g_i(A)$ for each $A \in \mathcal{F}$, $i \in [l]$, we will have $O(n \cdot l)$ sketches where $n = |\mathcal{F}|$. This is our LSH data structure.

For each value $i \in [l]$, a hash table should also be created to look up a set stored at index $g_i(Q)$ for some query set Q in constant time. Furthermore, each set $A \in \mathcal{F}$ stored in this hash table should have an associated hash table over its elements.

Creating the data structure takes $O(\sum_{A \in \mathcal{F}} |A| \cdot l)$ time and space during the pre-processing step.

To query, we first calculate l sketches of Q , $g_0(Q), \dots, g_{l-1}(Q)$, which takes $O(k \cdot l \cdot |Q|)$ time. These sketches can be used to retrieve l candidate sets $\mathcal{M} = \{A_0, \dots, A_{l-1}\}$ from each of the l outermost hash tables.

For each $A \in \mathcal{M}$, the Jaccard similarity can be computed in $O(|Q|)$ time by computing the size of the intersection using the associated hash table to A and computing the size of the union by the inclusion-exclusion principle if the size of A is known.

When an element $A \in \mathcal{M}$ is found such that $J(A, Q) \geq j_2$, the algorithm can return. This algorithm can therefore query in $O(l \cdot k \cdot |Q|) = O(n^\rho \log_2(n) |Q|)$ time, and since we use $l = \lceil n^\rho \rceil$ different hash functions combined with the $(j_1, j_2, 1/n^\rho, 1/n)$ -sensitivity of \mathcal{G} , the algorithm has a constant error probability [11].

2.5 Later Improvements

Many improvements have since been discovered. Andoni and Indyk [4] have since improved the total query time to $O(n^\rho |Q|)$ by using a technique called *tensoring*. Tensoring entails creating a collection of $m = o(l)$ sketch functions instead of the usual k by picking m functions from \mathcal{H} uniformly at random. Let each g_i for $i \in [l]$ from the LSH framework now be a distinct combination of these m hash functions of size t (where t is a positive integer). It is now possible to choose t such that $\binom{m}{t} = l$, such that every combination of the m hash values are used. By doing this, it is only necessary to compute $O(l)$ hash values of Q which is a reduction by a factor of k [4], [8].

Dahlgaard, Knudsen, and Thorup [7] have also introduced an improved sketch called the *Fast Similarity Sketch* (FSS), which is much faster to create and preserves the same properties as the MinHash sketch, while reducing the query time to $O(k \cdot l + |Q|) = O(n^\rho \log_2(n) + |Q|)$.

This improvement in query time is due to multiple factors.

First of all, a sketch of size k of a set A takes an expected running time of $O(k \log_2(k) + |A|)$ to create using the FSS algorithm, which is improved compared to the $O(k|A|)$ running time of the classic MinHash algorithm.

Secondly, the algorithm uses a single intermediate sketch that it samples the hash values of the sketch from, where the size of the intermediate sketch is independent of the size of the query set, which turns the time to create the sketch of Q from $O(n^\rho \log_2(n) |Q|)$ to $O(n^\rho \log_2(n) + |Q|)$.

Thirdly, the algorithm also uses a faster method for eliminating “bad matches” from the $O(l)$ candidate sets by estimating if the Jaccard similarity between the query set and candidate set is above some threshold within a certain probability. Here, a “bad match” is defined as a candidate set A where $J(A, Q) < j_2$. This reduces the filtering step to only take $O(l + |Q|)$ time, since the bad matches can be eliminated in $O(l)$ time and the amount of matches remaining is $O(1)$. These $O(1)$ matches can then finally be compared with to the query set which takes $O(|Q|)$ time.

These three factors result in a $O(n^\rho \log_2(n) + |Q|)$ total query time, which can be combined with tensoring to a $O(n^\rho + |Q|)$ query time as shown by Christiani [8].

Until now, the data structures presented error with a constant probability $O(1)$. If one wants to reach a better error probability ε (could for example be $\varepsilon = O(1/n)$), it is common practice within the field to use $M = \Theta(\log_2(1/\varepsilon))$ independent data structures. This will usually result in a query time of $O((n^\rho + |Q|) \log_2(1/\varepsilon))$.

Knudsen [11] has shown a method of achieving an improved $O(n^\rho \log_2(1/\varepsilon) + |Q|)$ query time instead, which is very beneficial when working with highly dimensional data sets and a small ε . This is done by sampling the $\Theta(M \cdot l)$ different sketches of Q necessary for querying from a single sketch of size $\Theta(M \cdot \log_2^2(n))$, with each sketch segment fed into a LSH data-structure.

Knudsen [11] has further improved the query time to $O((\frac{n \log_2(d)}{d})^\rho \log_2(1/\varepsilon) + |Q|)$ (where d is the word size) by reducing k such that \mathcal{H} becomes $(j_1, j_2, O((n/b)^\rho), n/b)$ -sensitive where $b = d / \log_2(d)$. Then, using a filtering algorithm inspired by [11], each of the l data structures will pick either a candidate set $A \in \mathcal{F}$ which is estimated to have $J(A, Q) \geq j_2$ or return none.

From the $\leq M$ different candidate sets provided by the M data structures, one needs to filter out those data structures that have a high probability of returning a “bad” set. If the amount of candidate sets are above a specific threshold, Knudsen [11] shows that

one can pick a random candidate set which will be correct with the required probability. If the amount of candidate sets is smaller than the threshold however, one needs to filter out the bad sets efficiently. This is done by keeping another sketch $P(A)$ for each $A \in \mathcal{F}$ of $\Theta(\max(\log_2(n), \log_2(1/\varepsilon)))$ hash values during the pre-processing step. For each of these hash values, only the least significant bit will be stored (also known as the b -bit minwise hashing trick [6]) as well as a set of $d = O(\log_2(w))$ coordinates of P . By comparing the sketch $P(A)$ with $P(Q)$ for each of the l candidate sets, we can eliminate bad matches with a probability that is high enough to achieve the desired error bounds. Since each sketch requires d bits, we can embed $O(w/\log_2(w))$ sketches into each word. When querying, instead of comparing our sketches of Q with each of the l candidate sketches one at a time, we can compare it to $O(\frac{d}{\log_2(d)})$ different sketches at once.

By calculating the amount of bits per sketch for A and Q that are equal, one can estimate the Jaccard Similarity to a degree which is good enough to filter a candidate match. If this algorithm runs in $O(n + \log_2(d))$ time, then Knudsen [11] shows that one can achieve the sub-linear $O((\frac{n \log_2(d)}{d})^\rho \log_2(1/\varepsilon) + |Q|)$ running time as claimed. This calculation can be done by bit-wise XNOR'ing a word packed with sketches of a set $A \in \mathcal{F}$ with a word packed with sketches of Q . Then, one can calculate the amount of bits set per sketch. Every bit set will then indicate that Q and A hashed to the same value, which means that we can filter out a group of matches if they belong to a word where none of the sketches contain an amount of bits set over some threshold. It is this calculation on how to count these bits efficiently that the rest of this project will focus on.

3 Parallel Bit Counting

In the same article where Knudsen [11] presents the improvements for algorithms with an $\varepsilon = o(1)$ error probability, a parallel algorithm to compute the cardinality of a sequence of bit-strings packed into sequence of words is also presented. It is only described by recurrences, and omits implementation details. I will try to describe the same algorithm in a much simpler fashion and prove its correctness and time complexity while doing it. This means that this section is my own interpretation of the algorithm, proofs and analysis described in [11].

For a word t with word size w packed with w/d bit-strings for some $d \leq w$, a naive algorithm to perform bit-counting could be described like in algorithm 1.

Algorithm 1 A naive linear time algorithm

```

1: function CARDINALITY( $t$ )
2:    $S \leftarrow \{\}$ 
3:   for  $j \in [w/d]$  do
4:      $x \leftarrow 0$ 
5:     for  $i \in [d]$  do
6:        $x \leftarrow x + ((t \ll j \cdot d) \gg i) \wedge 1$ 
7:     end for
8:      $S \leftarrow S \cup \{x\}$ 
9:   end for
10:  return  $x$ 
11: end function

```

This algorithm trivially runs in linear time to the word size¹ w of the input word t . For n words of size w , this algorithm runs in $O(nd \frac{w}{d}) = O(nw)$ time. Knudsen [11] presents two improvements to this: The first will improve the running time to $O(n \log(d))$ by utilizing a divide-and-conquer technique and the second to $O(n + \log(d))$ time by introducing parallelism.

3.1 Divide-and-Conquer

To explain how divide-and-conquer methods can be used to improve running time, we must first introduce a word sized bit mask from [11] that is defined like so (in binary):

$$m_{i,j} = \underbrace{0 \dots 0}_{2^j - 2^i} \underbrace{1 \dots 1}_{2^i} \dots \underbrace{0 \dots 0}_{2^j - 2^i} \underbrace{1 \dots 1}_{2^i}$$

where $j > i$ and j and i are non-negative integers. The notation m_i is a shorthand for $m_{i,i+1}$ and indicates a mask with an equal amount of 1's and 0's. By computing $t \wedge m_{i,j}$ for some word t and some integers i, j , we can replace every other sub-string of size $2^j - 2^i$ in the bit-string with 0's. This will be described as *isolating* segments of size 2^i .

We will use this in the following operation:

$$T(t, m, k) = (t \gg k) \wedge m \tag{1}$$

¹For the rest of this project, it is assumed that w is a power of 2.

This operation isolates the segments indicated by the bit-mask starting from the k th position of the word t . This is useful for partitioning a bit-string, since any word now can be described as

$$t = T(t, m_i, 0) + (T(t, m_i, 2^i) \ll 2^i)$$

Notice that even though we use addition, overflows are impossible since the segment pattern of $T(t, m_i, 0)$ and $(T(t, m_i, 2^i) \ll 2^i)$ do not overlap.

The algorithm to calculate the cardinality of a single bit-string in a word t can then be described like in algorithm 2

Algorithm 2 A divide-and-conquer approach

```

1: function CARDINALITY( $t$ )
2:   for  $i \in [\log_2(d)]$  do
3:      $t \leftarrow T(t, m_i, 0) + T(t, m_i, 2^i)$ 
4:   end for
5:   return  $t$ 
6: end function

```

The simplest way to gain intuition of this algorithm is to prove it. A loop-invariant is presented like so:

Invariant 1. *At the beginning of the i th iteration of the algorithm, the word $t^{(i)}$ can be regarded as consisting of $w/2^i$ bit-string-segments of size 2^i , that represents the cardinality of the corresponding segments of the word $t^{(0)}$.*

Proof. Base case: When $i = 0$, then each bit-string of size $2^0 = 1$ in $t^{(0)} = t$ will be 1 if the bit is 1 and 0 if the bit is 0. This proves our base case.

Induction step: At step i , invariant 1 states that the word $t^{(i)}$ contains $\frac{d}{2^i}$ segments of size 2^i . The operation $T(t, m_i, 0)$ will then isolate exactly every other segment of size 2^i . The operation $T(t, m_i, 2^i)$ will isolate the remaining segments and bit-shift them such that the segments align with the segments isolated by $T(t, m_i, 0)$.

By adding $T(t, m_i, 0)$ and $T(t, m_i, 2^i)$ together, every pair of overlapping segments will be added together to form a new segment of double the size. Invariant 1 states that each segment of $T(t, m_i, 0)$ and $T(t, m_i, 2^i) \ll 2^i$ represents the cardinality of a corresponding segment in $t^{(0)}$. Since the cardinality of a bit-string is equal to the sum of its parts, we can interpret each new segment as the cardinality of the corresponding two previous segments combined. Therefore, at the end of the i th iteration of the for-loop, the word t can be interpreted as consisting of bit-string segments of size 2^{i+1} , where each bit-string-segment represents the cardinality of the corresponding bit-string segments of $t^{(0)}$. \square

When the algorithm terminates after $\log_2(d)$ iterations, the segments will have size $2^{\log_2(d)} = d$ and thus span the entire size of the original bit-string, which means that we have the bit-count of the original bit-string. If m_i for $i \in [\log_2(d)]$ is pre-computed, then each iteration can be performed in constant time. If the word t instead contains a sequence of $m = w/d$ bit-strings of size d , the cardinality of all of the m bit-strings can be calculated in $O(\log_2(d))$ time. If one has n of such words, the running time will therefore be $O(n \log_2(d))$. It is to be noted that this algorithm returns a word $t^{(\log_2(d))}$ where the cardinalities of each of the bit-strings from $t^{(0)}$ are embedded inside the word.

3.1.1 Example

Given a word $t^{(0)} = 1101011001010010_2$ containing 1 bit-string, we can run the algorithm by hand to find the cardinality. Due to invariant 1, we can regard $t^{(0)}$ as consisting of $w/2^0 = 16$ sections of size $2^0 = 1$ as marked by colors:

$$t^{(0)} = \textcolor{blue}{1}\textcolor{red}{1}\textcolor{blue}{0}\textcolor{red}{1}\textcolor{blue}{0}\textcolor{red}{1}\textcolor{blue}{1}\textcolor{red}{0}\textcolor{blue}{0}\textcolor{red}{1}\textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{1}\textcolor{red}{0}_2$$

At the beginning of the first iteration, we can calculate m_0 to be 0101010101010101_2 . This means that

$$T(t^{(0)}, m_0, 0) = 0101010001010000_2$$

and

$$T(t^{(0)}, m_0, 2^0) = 0100000100000001_2$$

We can add them together to get $t^{(1)}$:

$$t^{(1)} = \textcolor{blue}{1}\textcolor{red}{0}\textcolor{blue}{0}\textcolor{red}{1}\textcolor{blue}{0}\textcolor{red}{1}\textcolor{blue}{0}\textcolor{red}{1}\textcolor{blue}{0}\textcolor{red}{1}\textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{0}\textcolor{red}{1}_2$$

Notice that each section has now doubled in size. Due to invariant 1, each section represents the bit-string cardinality of the corresponding section of $t^{(0)}$. For example, the least significant section $\textcolor{red}{01}_2 = 1_{10}$ represents that there is 1 bit set in the 2 least significant bits of $t^{(0)}$ (which we can see is true). Continuing the algorithm, we get:

$$t^{(2)} = \textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{1}\textcolor{red}{1}\textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{1}\textcolor{red}{0}\textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{1}\textcolor{red}{0}\textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{0}\textcolor{red}{1}_2$$

$$t^{(3)} = \textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{0}\textcolor{red}{1}\textcolor{blue}{0}\textcolor{red}{1}\textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{0}\textcolor{red}{0}\textcolor{blue}{0}\textcolor{red}{1}\textcolor{blue}{1}_2$$

$$t^{(4)} = \textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{0}\textcolor{red}{1}\textcolor{red}{0}\textcolor{red}{0}_2 = 8_{10}$$

The algorithm has now terminated, and we can see that $t^{(0)}$ contains 8 bits set, which we can see is true. If instead, we regarded the word $t^{(0)}$ as containing two bit-strings of size 8 each, then we could stop the algorithm after iteration $i = 2$, where $t^{(3)}$ contains the bit-string cardinalities of each of the two sections. Then we could see that the first bit-string contains 5 bits set, and the second contains 3 bits set.

3.2 Parallelism

To introduce parallelism into the algorithm, we must first realize the following: When two segments of 2^i bits gets combined, they will not need all of the 2^{i+1} bits to represent their sum. It is actually such that the bits used at iteration i is exactly $i + 1$.

Invariant 2. *At the end of the i th iteration of algorithm 2, the amount of bits set in a given segment of a word $t^{(i+1)}$ is at most $i + 1$.*

Proof. At iteration i , the algorithm will add two words where each bit-string-segment has a size of 2^i bits. In the worst case, where all bits are set, each segment in $t^{(i)}$ will have the value $2^{i+1} - 1$, which means that each segment in $t^{(i+1)}$ will have the value $2^{i+1} - 1 + 2^{i+1} - 1 = 2(2^{i+1} - 1)$. To find out how many bits is needed to represent the given number, we can calculate \log_2 of the expression and round down: $\lfloor \log_2(2(2^{i+1} - 1)) \rfloor = \lfloor 1 + \log_2(2^{i+1} - 1) \rfloor = 1 + \lfloor \log_2(2^{i+1} - 1) \rfloor = i + 1$. Therefore, at the end of the i th iteration, each segment can only have $i + 1$ or fewer bits set. \square

Now, we will introduce a function $l(i)$, which produces the smallest non-negative integer such that $2^{l(i)} \geq i + 2$. We can now use the bit-mask $m_{l(i), i+1}$ instead of m_i , since the segments are guaranteed to be strictly smaller than $2^{l(i)}$. At the end of each iteration, there will be a sequence of $2^{(i+1)-l(i+1)}$ unused bits between each segment, since each segment will have size $2^{(i+1)}$ (invariant 1) but at most use $i + 1$ bits (invariant 2). In the case that $2^{(i+1)-l(i+1)} \geq 2^{l(i+1)}$, this unused sequence can contain a segment from an entirely different word (note that this happens when $l(i) = l(i+1)$). Algorithm 3 uses this to efficiently compute the cardinality of all the words in a set $S = \{t_0, \dots, t_{n-1}\}$.

Algorithm 3 A parallel divide-and-conquer algorithm

```

1: function COMPUTE( $S, d$ )                                 $\triangleright S$  is the input set,  $d$  is the word size
2:    $n \leftarrow S.length$ 
3:   for  $i \in [\log_2(d)]$  do
4:      $k' \leftarrow \{\}$ 
5:      $t \leftarrow \{\}$ 
6:     for  $t \in S$  do
7:        $k' \leftarrow k' \cup \{T(t, m_i, 0) + T(t, m_i, 2^i)\}$ 
8:     end for
9:     if  $i < 2$  then
10:       $S \leftarrow k'$ 
11:    else
12:      if  $l(i) = l(i+1)$  AND  $S.length > 1$  then
13:        for  $j \in \lceil [S.length/2] \rceil$  do
14:           $t \leftarrow t \cup \{k'_{2j} + (k'_{2j+1} \ll 2^i)\}$ 
15:        end for
16:      else
17:        for  $j \in k'$  do
18:           $t \leftarrow t \cup \{T(k'_j, m_{l(i)}, 0) + (T(k'_j, m_{l(i)}, 2^{l(i)}) \ll 2^i)\}$ 
19:        end for
20:      end if
21:       $S \leftarrow t$ 
22:    end if
23:  end for
24:   $S' \leftarrow \{\}$ 
25:  for  $j$  in  $[S.length]$  do                                 $\triangleright$  For every word in the final set
26:    for  $k$  in  $[2^{l(\log_2(d))}]$  do                         $\triangleright$  For every original bit-string embedded in  $S_j$ 
27:       $S' \leftarrow S' \cup \{T(S_j, m_{l(\log_2(d)), \log_2(d)+1}, k \cdot 2^{l(\log_2(d))})\}$ 
28:    end for
29:  end for
30:  return  $S'$ 
31: end function

```

This algorithm is quite a bit more complex than the first one, and to prove it we need some more terminology.

When we pack the words $t_0, \dots, t_{2^{i-l(i)}-1}$ into a word t , we say that t is i -packed. This means one can extract all $2^{i-l(i)}$ different words by performing the operation:

$$t_j^{(i)} = T(t, m_{l(i), i+1}, j \cdot 2^{l(i)}) \quad (2)$$

for all $j \in [2^{i-l(i)} - 1]$. The point of the algorithm is to ensure that after every step of the algorithm that every word in S is $(i + 1)$ packed. When the first loop of the algorithm terminates at $i = \log_2(d) - 1$, then every word will be $\log_2(d)$ -packed, which means that it fits $2^{\log_2(d) - l(\log_2(d))}$ words. When working with 64-bit words, this will result in $2^{6-3} = 8$ words per packed word.

Invariant 3. *At the end of the i th iteration of the first loop of the algorithm, every word in S will be $(i + 1)$ -packed for any $i \geq 1$*

Proof. First, the first two iterations of the loop are run, such that $i = 1$. During the first two iterations of the loop, no words are combined since it just uses the naive algorithm, which upholds the invariant because any 2-packed words embeds $2^0 = 1$ of the original words.

For any iteration where $i \geq 2$, one out of two things can happen. Either, $l(i + 1) = l(i)$ (we do not need more bits to describe each segment), or $l(i + 1) = l(i) + 1$ (we need twice the amount of bits to describe each segment).

In the first case, we will combine the words using bit-shifting. Since every word in k' uses the same amount of bits to represent each segment even though we just combined segments when creating k' , there must be an empty space of at least size $2^{l(i+1)}$ in every segment. We can fill out that empty space by adding another word bit-shifted by a factor of 2. Now, the amount of words embedded is equal to the sum of both of its parts (i.e. has doubled). This upholds the loop invariant since any $(i + 1)$ -packed word must have twice the amount of words embedded than a i -packed word if $l(i + 1) = l(i)$ by definition.

In the second case, $l(i + 1) = l(i) + 1$. In this case, we need to relocate our segments that actually use $2^{l(i)}$ bits so they fit into the segments of size 2^i bits that were created when creating k' . The same amount of words have been embedded into each words as before, but as $l(i + 1) = l(i) + 1$ a $(i + 1)$ -packed word should only embed the same amount of words as a i -packed word, which upholds the invariant. Furthermore, since we relocated the segments, we can now still extract each original word by performing the operation in equation 2.

The for-loop will terminate when $i = \log_2(d) - 1$, from which each word in $S^{(\log_2(d))}$ will embed $2^{\log_2(d) - l(\log_2(d))}$ of the words from $S^{(0)}$. \square

When the first for loop has finished, each S'_i for $i \in [n]$ will contain words with embedded cardinalities due to invariant 1, just like algorithm 2.

3.2.1 Example

For the first two iterations of the first for-loop, the algorithm works exactly as algorithm 2. When $i = 2$, lines 17 – 19 are hit for the first time, but since the words are (3)-packed, there will only be $2^{2-2} = 1$ words from S^0 embedded in each word in $S^{(2)}$. This results in the operation at line 18 leaving k' unchanged.

When $i = 3$, the lines 12 – 16 are hit for the first time. This results in words getting pairwise combined. Here is an example of how words in k' for some set $S^{(3)}$ can get merged together (segments of size $2^{l(i)}$ are colored):

$$k' = \{00000000\textcolor{red}{00001111}00000000\textcolor{red}{00001001}_2, 00000000\textcolor{blue}{00001010}00000000\textcolor{blue}{00010000}_2\}$$

3.4 Time Complexity Analysis

We can now move on to showing the time complexity of the algorithm. First, we can try to find the running time of the first for-loop. At each iteration i , we will describe the amount of elements in $S^{(i)}$ as $n^{(i)}$ with $n = n^{(0)}$. We can then describe $n^{(i)}$ like so:

$$n^{(i)} = \begin{cases} n & \text{if } i \leq 1 \\ \lceil \frac{n}{2^{i-l(i)}} \rceil & \text{otherwise} \end{cases} \quad (3)$$

which means that the first for-loop in total takes

$$\begin{aligned} & \sum_{i=0}^{\log_2(d)} n^{(i)} \\ &= 2n + \sum_{i=2}^{\log_2(d)} n^{(i)} \\ &= 2n + \sum_{i=2}^{\log_2(d)} \lceil \frac{n}{2^{i-l(i)}} \rceil \end{aligned}$$

Since the last loop iterates over every element of the original set, it must take $O(n)$ time. We can now derive the total running time:

$$\begin{aligned} & O(n) + 2n + \sum_{i=2}^{\log_2(d)} \lceil \frac{n}{2^{i-l(i)}} \rceil \\ &= O(n) + \sum_{i=2}^{\log_2(d)} \lceil n2^{l(i)-i} \rceil \\ &\leq O(n) + \sum_{i=2}^{\log_2(d)} (n2^{l(i)-i} + 1) \\ &\leq O(n + \log(d)) + n \sum_{i=2}^{\log_2(d)} 2^{l(i)-i} \end{aligned}$$

Now we use the fact that $l(i)$ is always the smallest number that satisfies $2^{l(i)} \geq i + 2$. Since it is the smallest number, then $2^{l(i)} < 2(i + 2)$, which can be proved by contradiction: $2^{l(i)} \geq 2(i + 2) \implies 2^{l(i)-1} \geq i + 2$, which shows that there is a number that is 1 smaller than $l(i)$ that upholds the bound $2^{l(i)} \geq i + 2$. Therefore $2^{l(i)} < 2(i + 2)$. We can now use this fact:

$$\leq O(n + \log(d)) + n \sum_{i=0}^{\infty} \frac{2(i + 2)}{2^i} = O(n + \log(d))$$

The running time of the algorithm can therefore be bound to $O(n + \log(d))$.

4 Empirical Evaluation

4.1 Methods

For the rest of the project, the main focus will be on showing the practical feasibility of the parallel bit-counting algorithm compared to simpler algorithms. Since d needs to be very large for parallelism to affect the results, it can be very hard to compare the algorithms presented in the previous section in a meaningful way if multiple bit-strings should be embedded in any word sizes found on a standard machine. For the rest of the project, it therefore be assumed that $d = w$, which means that only 1 bit-string will be embedded per word.

We know from the time complexity of the algorithms in the previous section that the parallel algorithm gets theoretically better than the other algorithms as d grows. If we find out that the parallel algorithm is slower than the other algorithms in the case that $d = w$, then we can assume that it is also slower when $d \leq w$. If we however find out that the parallel algorithm is faster, then we might be unable to make any conclusions on the algorithms feasibility, since its usefulness in the filtering step of the querying algorithm relies on its ability to work on words with multiple embedded bit-strings.

A comparison will be made to the `popcnt` CPU instruction that is implemented on most x64-86 CPUs, which can count the bit-wise cardinality of a word [16].

4.1.1 Benchmarking

To be able to quantify whether one algorithm is faster than the other, one can try to implement it in practice, run it on a specified data set and measure how long it takes for the algorithm to terminate. Benchmarking this way is common in the industry, but it can be hard to be able to generalize from results, due to the many possible sources of error such as:

- Results can vary from machine to machine due to factors like difference in instruction sets, processor infrastructure, cache hierarchy, compiler compatibility, memory capacity etc.
- The input data set can be biased towards one specific kind of architecture or implementation.
- Caches become more efficient when “warmed up” e.g. when the computer has recently fetched from the same memory addresses. Two identical benchmarks might produce different results based on whether or not the caches are filled.
- A process that benchmarks a program will have to share the CPU cores with other processes running on the computer. This can produce a high variance between identical benchmarks on the same machine.
- Different compilers can produce more or less optimized code, which might produce varying results on different compilation targets.

The scope of this project disallows running benchmarks on a wide variety of computers, which presents some uncertainty in the validity of the results. Some of the other sources of error are mitigable.

- The benchmark should be run multiple times with different pseudo-randomly generated input to mitigate bias in the input data.
- The benchmark should be run multiple times per input data set, until a significant mean execution time can be calculated.
- The benchmark should perform a few "dry runs" before each run to warm up the caches.

This applications of this software lends itself to large-scale machine learning applications, search engines and data mining. Therefore, the most relevant results would come from running the benchmarks on a machine with access to a GPU, and for the algorithm to be written with parallelism² in mind. This has not been considered within the scope of this project however, but instead lends itself as an open research problem.

4.2 Implementation

4.2.1 Technology

The implementation to be measured upon has been written in Rust, a memory safe, fast, compiled programming language [14] that has great tooling and typing for writing, testing and benchmarking algorithms.

It uses the `rustc` compiler with a LLVM compilation back-end. It does not use a garbage collector but relies on compile time borrow checking, which leads to a language that is comparable to or faster than C [10].

The binary to be benchmarked is compiled and run on a laptop workstation as described in appendix A. The code for the implementation is available in appendix B.

4.2.2 Design

One of the main design challenges of the implementation is to support multiple word sizes simultaneously to give an accurate estimate on how the execution time scales with the word size. This has the consequence of making it harder to hard-code values like $m_{i,j}$ and $l(i)$, since a theoretical implementation with a very large word size would have to compute these values at runtime.

This implementation uses hard-coded values however, since any real-life implementation with a reasonable word size would do so as well.

Another challenge is code optimization. Rust has an `#[inline(always)]` macro that forces function inlining, which reduces the time needed to handle the stack. A nice feature about the algorithm is that it does not use multiplication, which usually is a very slow operation. It does however operate on very large lists of numbers, which need to be handled efficiently. For this, the heap-allocated `Vec` data structure was used, as it functions as a dynamic array and contains many features to help with optimizations. For instance, the `Vec::with_capacity(n)` constructor specifies an initial capacity of the dynamic array such that unnecessary allocations are avoided, and the `Vec::truncate()` allows constant time shortening of arrays [14].

²In this case, parallelism is meant in a GPU context unlike word parallelism as mentioned previously.

4.2.3 Correctness

It is luckily easy to verify the correctness of the implementation of the algorithm since it is a deterministic algorithm. That means that it is possible to calculate any expected output by hand, which both is useful for validating and for debugging.

To be able to show that the implementation works on a wide array of inputs, we can generate a large list of random numbers to help cover eventual edge cases. We can then compare the results to some reference implementation - in this case the built-in `count_ones()` method.

It is much harder to verify whether the code works with multiple different word sizes. This is due to the fact that the else branch on line 16 of algorithm 3 only gets hit once when working with $w = 64$ (when $i = 2$) and twice when $w = 128$. Since the applications of this algorithm really benefits when using large word sizes, it is very relevant that this gets tested.

4.2.4 Benchmarking

To benchmark the algorithm, two reference implementations were created to be used as a baseline models alongside the CPU instruction. These implementations were based on algorithm 1 and 2. A list of 2^{24} random numbers were to be generated and used as input for each of the algorithms.

The Rust library `Criterion.rs` [12] was used to perform the benchmarking. `Criterion.rs` was very useful for multiple reasons.

- First of all, `Criterion.rs` integrates into standard Rust tooling easily and allows benchmarking on binary code built for production (e.g. with full optimizations and no debugging symbols)
- Secondly, `Criterion.rs` provides a compiler "opaque" functions that makes sure that repeated calls to the same function do not get optimized away.
- `Criterion.rs` handles warming up caches by running the algorithms multiple times before beginning to measure.
- Lastly, `Criterion.rs` does not record the execution time of a single operation - instead it runs multiple samples each consisting of a number of iterations. The running time is then calculated as the execution time per sample divided by the amount of iterations.

4.3 Results

The benchmarks were run on different input sizes that were all powers of two - in this case $\{2^4, 2^5, \dots, 2^{19}\}$ and with word sizes $\{16, 32, 64, 128\}$. The results can be seen in figure 1. It can be seen that the parallel algorithm is significantly slower for small word sizes, and only becomes more efficient than the naïve sequential algorithm when $w = 128$. This must mean that the hidden factors hidden in the big-O notation of the parallel algorithm's time complexity outweigh the w factor of the naïve algorithm. This makes sense, since the parallel algorithm is much more complex and only starts benefiting from parallelism after the first three iterations. Interestingly enough, when $w = 64$, the parallel algorithm only seems to be slower when the input size is larger than 2^{16} . This can be due to the input data being able to fit inside 128 KiB L1 data

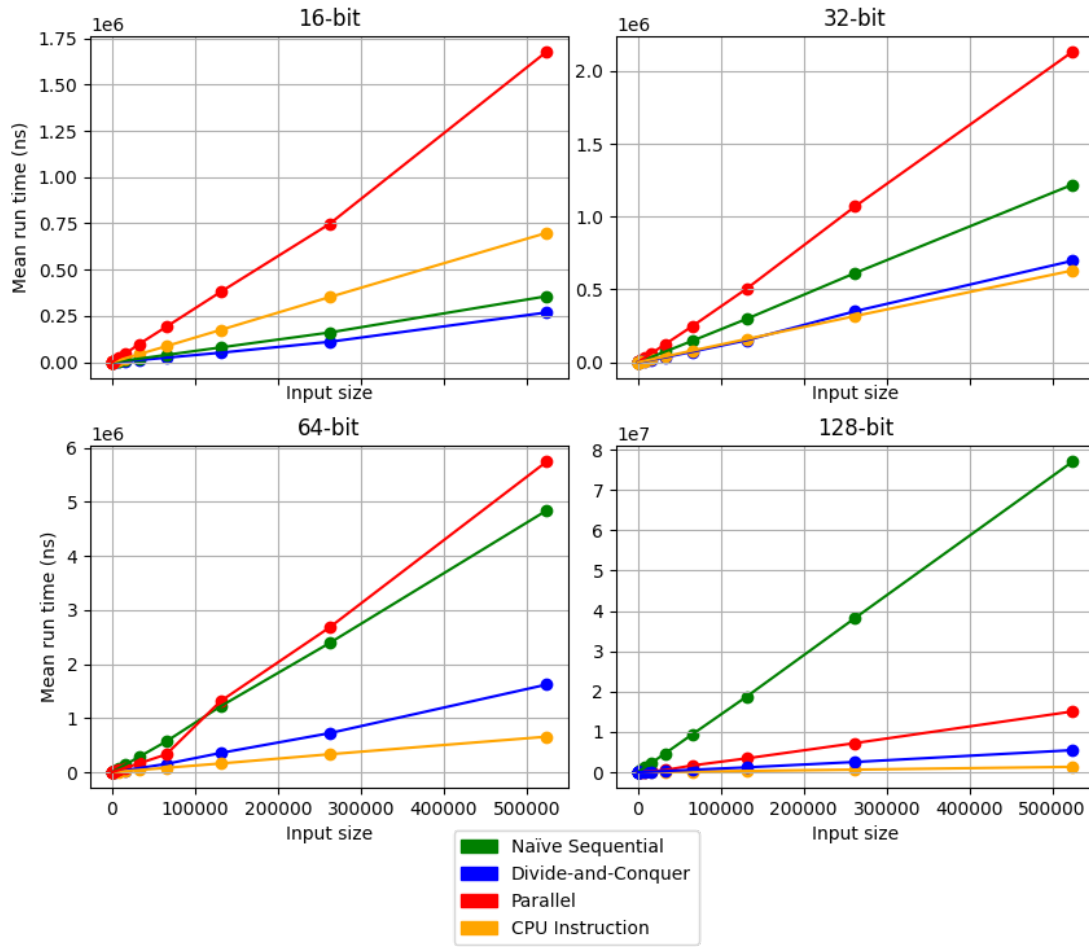


Figure 1: Execution time of each algorithm on different word sizes

cache on my machine (see appendix A for machine specifications).

When run on 128-bit size words, the naïve sequential algorithm performs much worse than both the parallel and the divide-and-conquer algorithm. This shows that when we increase the word size, algorithms that have a linear dependency on w suffer much more than algorithms that have a logarithmic dependency, which is as expected.

It is also seen that the divide and conquer algorithm significantly outperforms both the parallel and the naïve algorithm at all word sizes. This might be due to the fact that it is much simpler than the parallel algorithm and is asymptotically more optimal than the naïve sequential algorithm. The final thing to note is that the `popcnt` CPU instruction is the ultimately fastest option out of the algorithms. This shows that these algorithms might only be useful on hardware that does not support this instruction or on very large word sizes. Even then, the divide-and-conquer algorithm seems to be sufficiently efficient for many uses or even superior to the parallel algorithm despite its sub-optimal theoretical running time.

4.4 Discussion

Even though the results indicate a clear advantage to simpler algorithms to perform bit counting, the parallel bit counting algorithm should not be completely dismissed. While we can be pretty sure that a smaller $d < w$ might not improve the performance of the parallel bit-counting algorithm compared to the other algorithms, we do not know how the parallel algorithm performs on very large word sizes (e.g. 256 to 384 bit words) as seen on some GPU units [13]. Since the applications of the similarity search problem usually apply to very large data sets, it is often appropriate to perform highly efficient parallel calculations when possible. Parallel programming comes with many new challenges that were not considered in this project such as minimizing read/write operations to the main memory of the computer, the memory capacity of the graphics card and the size of the register files. This makes an actual implementation of the algorithm 3 non-trivial, especially since most programming interfaces for interacting with a GPU relies on intricate low-level systems programming (unless one were to use a language like Futhark [15]). A nice feature of the algorithm is that it does not use multiplication, which means that many of the computations might be efficiently computed on a GPU.

Another thing that this experiment does not take into account is the index calculation time required to actually look up into the list of cardinalities. This was not included, since the correctness criteria of the algorithm may or may not take into account whether or not the cardinalities should have the same index in the final list as the input data.

The implementation that this experiment relied upon was optimized as per the best of my abilities, but further optimizations might be possible. A nice feature of the Rust programming language is that the time complexity is very often documented in the documentation, which was very useful when working with heap-allocated vectors of numbers. The implementation was also written to be as generic as possible in regards to word size such that the experiment would be performed as unbiased as possible. This was possible due to Rust's generics, and while Rust design philosophy relies upon zero-cost abstractions [14] (like C++), this might cause some optimization issues. Further optimizations might also be possible by manually analyzing the produced as-

sembly code, which was regarded as out-of-scope for this project. In general, while parallel bit counting theoretically is more efficient than its sequential and instruction-level counterparts, it has yet to be shown to be the case in practice. While further research might research how it performs on specialized hardware, the algorithm behind this will probably first be useful when larger word sizes become more easily available and computable, if ever.

5 Conclusion

An in-depth analysis of existing solutions to the Approximate Similarity Search Problem described how many existing solutions utilize sketching to achieve a low query time with a constant error probability. In the case that one wants a sub-constant error probability, a possible solution by Knudsen [11] was presented that can reduce the query time even further.

An important part of this optimization relies in being able to quantify the accuracy of a collection of data structures very efficiently, which is done in part by calculating the cardinality of a sequence of bit-strings packed into a list of words. Knudsen [11] presents an algorithm to do this in parallel, which reduces the total query time to be sub-linear to the amount of sets in the corpus.

An analysis of the bit-counting algorithm shows its correctness with a few modifications alongside a theoretical running time which matches [11]. Benchmarks performed on an actual implementation tells otherwise however, showing that the word parallelism used to achieve a low amortized cost actually is hard to match simpler algorithms in practice when used on standard word sizes. Further research might focus on benchmarking the algorithm on more specialized hardware, although rewriting the implementation seems to be non-trivial.

References

- [1] A. Broder, *On the resemblance and containment of documents*, 1997. DOI: 10.1109/SEQUEN.1997.666900.
- [2] P. Indyk and R. Motwani, “Approximate nearest neighbors: Towards removing the curse of dimensionality,” in *STOC ’98*, 1998.
- [3] A. Gionis, P. Indyk, and R. Motwani, “Similarity search in high dimensions via hashing,” *Proceeding VLDB ’99 Proceedings of the 25th International Conference on Very Large Data Bases*, vol. 99, May 2000.
- [4] A. Andoni and P. Indyk, “Efficient algorithms for substring near neighbor problem,” Jan. 2006, pp. 1203–1212. DOI: 10.1145/1109557.1109690.
- [5] P. Li and A. C. König, “Theory and applications of b-bit minwise hashing,” *Commun. ACM*, vol. 54, no. 8, pp. 101–109, Aug. 2011, ISSN: 0001-0782. DOI: 10.1145/1978542.1978566. [Online]. Available: <https://doi.org/10.1145/1978542.1978566>.
- [6] P. Li, A. Shrivastava, J. Moore, and A. C. König, *Hashing algorithms for large-scale learning*, 2011. arXiv: 1106.0967 [stat.ML].
- [7] S. Dahlgaard, M. B. T. Knudsen, and M. Thorup, *Fast similarity sketching*, 2017. arXiv: 1704.04370 [cs.DS].
- [8] T. Christiani, *Fast locality-sensitive hashing frameworks for approximate near neighbor search*, 2018. arXiv: 1708.07586 [cs.DS].
- [9] S. Vassilvitskii, *Coms 6998-12: Dealing with massive data*, 2018.
- [10] P. Emmerich, S. Ellmann, F. Bonk, *et al.*, “The Case for Writing Network Drivers in High-Level Programming Languages,” in *ACM/IEEE Symposium on Architectures for Networking and Communications Systems (ANCS 2019)*, *jsstrong’s Best Paper Award*, Sep. 2019.
- [11] J. B. T. Knudsen, *Fast similarity search with low error probability*, 2021.
- [12] *Criterion.rs*, <https://github.com/bheisler/criterion.rs>, Accessed: 2021-12-09.
- [13] *Nvidia geforce rtx 3090 specs*, <https://www.techpowerup.com/gpu-specs/geforce-rtx-3090.c3622>, Accessed: 2021-12-22.
- [14] *Rust programming language*, <https://www.rust-lang.org/>, Accessed: 2021-12-06.
- [15] *The futhark programming language*, <https://futhark-lang.org/>, Accessed: 2021-12-22.
- [16] *The popcnt instruction*, <https://www.felixcloutier.com/x86/popcnt>, Accessed: 2021-01-05.

Mathematical Notation

Although sets by definition are not ordered, it is commonly seen in the literature to index into a set (see [11] for an example). Therefore, sets are regarded as ordered for this project.

Expressions within square brackets are indicator variables denoted using Iverson notation:

$$\begin{aligned}[X = Y] &\in \{0, 1\} \\ 0 &\leq Pr[[X = Y] = 1] \leq 1\end{aligned}$$

Integers within square brackets denotes the integer interval between 0 and the number.

$$[n] = [0 \dots (n - 1)] = \{0, 1, \dots, n - 1\}$$

Numbers in a different base than 10 will be indicated with a subscript.

$$11010110_2 = 214$$

Bit-shifting is done using C-style notation. Therefore

$$0101_2 \ll 1 = 1010_2$$

$$1010_2 \gg 1 = 0101_2$$

Logical operations between integers are meant as bit-wise.

$$\text{AND: } 0101_2 \wedge 0011_2 = 0001_2$$

$$\text{XNOR: } 0101_2 \odot 0011_2 = 1001_2$$

For this project, it is often useful to reference the same element after different iterations of an algorithm. To do this, the notation $x^{(i)}$ means an element after $i - 1$ iterations of an algorithm.

Appendices

A Machine Specifications

Model	Lenovo Thinkpad E580
Operating System	Arch Linux kernel 5.15.5-arch1-1
Word size	64bit
Processor model	Intel(R) Core(TM) i7-8550U CPU @ 1.80GHz
ISA	x86_64
Memory Capacity	32 GB
Caches	128 KiB L1d, 128 KiB L1i, 1MiB L2, 8 MiB L3

B Code

The code for the implementation part can be found on in this Github repository:
<https://github.com/Snailed/parallel-bit-counting>.